

Analysis Report: Model-Free vs Model-Based Reinforcement Learning

This report analyzes the behavioral difference between Model-Free and Model-Based Reinforcement Learning using a simple two-stage environment. The task demonstrates how standard Q-Learning can incorrectly assign credit to an action after a rare rewarded transition, while a Model-Based agent correctly reasons about the environment structure.

1. Environment Structure

The environment contains one initial state **S0** and two possible actions: **A1** and **A2**. The transition dynamics are probabilistic.

Current State	Action	Next State Probabilities
S0	A1	0.8 → S1, 0.2 → S2
S0	A2	0.2 → S1, 0.8 → S2

Reward probabilities: State **S1** gives reward with probability **0.8** — State **S2** gives reward with probability **0.2**

Therefore, the optimal long-term strategy is to choose **A1** because it usually leads to the more rewarding state **S1**.

2. Mathematical Blindspot of Q-Learning

The standard Bellman update equation for Q-Learning is:

$$Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$

where:

- alpha is the learning rate
- gamma is the discount factor
- r is the received reward

The Model-Free Q-Learning agent only learns from experienced trajectories. It does not explicitly understand transition probabilities or the causal structure of the environment. Suppose a rare event occurs: **S0** → **A2** → **S1** → Reward. Even though **A2** usually leads to **S2**, the received reward is propagated backward to **Q(S0,A2)**. The agent incorrectly increases preference for **A2** because the Bellman update only associates actions with obtained rewards.

The problematic term is: $r + \gamma \max_{a'} Q(s',a')$. This term directly credits the previously selected action for the observed reward without considering whether the transition itself was rare. Thus, the Model-Free agent is mathematically blind to the hidden structure: Action → State → Reward. Instead, it

learns only: Action $\rightarrow$ Reward

3. Model-Based Reinforcement Learning

A Model-Based agent explicitly stores a model of the environment. It must store:

1. Transition probabilities: $P(s'|s,a)$
2. Reward function: $R(s,a,s')$
3. State values: $V(s)$

The Model-Based action-value equation is:

$$Q_{MB}(s,a) = \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma V(s')]$$

For the given environment:

$$Q_{MB}(S0,A1) = 0.8*V(S1) + 0.2*V(S2)$$

$$Q_{MB}(S0,A2) = 0.2*V(S1) + 0.8*V(S2)$$

Since $V(S1) > V(S2)$ (reward probability 0.8 vs 0.2), the Model-Based agent correctly concludes that A1 is the better action because it more reliably reaches the high-reward state S1.

4. Hybrid Biological System

Research in neuroscience suggests that biological brains combine both habitual (Model-Free) and planning-based (Model-Based) systems. A simple hybrid formulation is:

$$Q_{Hybrid}(s,a) = (1 - \lambda) * Q_{MF}(s,a) + \lambda * Q_{MB}(s,a)$$

where:

- Q_{MF} is the Model-Free value
- Q_{MB} is the Model-Based value
- λ controls the balance between the two systems

Interpretation:

- $\lambda = 0 \rightarrow$ purely habitual behavior
- $\lambda = 1 \rightarrow$ purely planning-based behavior
- $0 < \lambda < 1 \rightarrow$ hybrid biological decision making

5. Conclusion

This experiment demonstrates the conceptual limitation of standard Q-Learning. A Model-Free agent incorrectly assigns credit to actions after rare rewarding outcomes because it lacks an internal model of the environment. In contrast, a Model-Based agent reasons about transition probabilities and future

state outcomes, allowing it to make more rational decisions. Biological systems likely combine both approaches, balancing fast habitual learning with flexible planning-based reasoning.